

# 03-credit-risk-model-feature-consumption

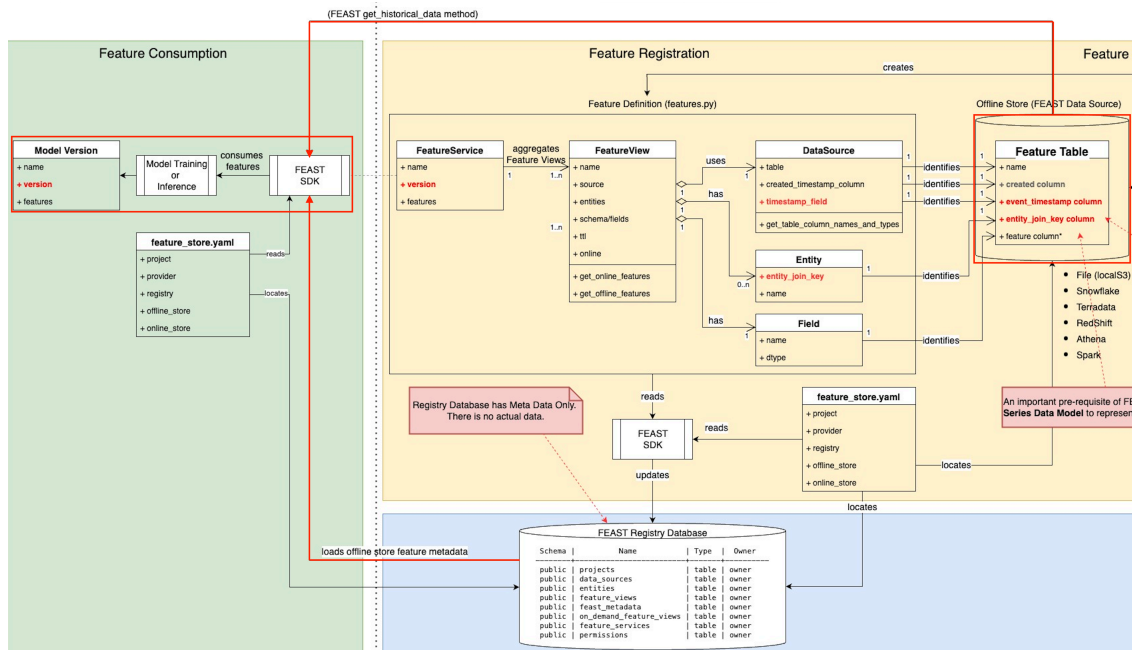
September 22, 2025

## 1 Feature Retrieval and Consumption for Model Training

The goal is to **Consume Features from the Offline Store** for model training.

### 1.1 This Notebook (03-feature-consumption)

This notebook goes through the steps of: 1. Load the Feature Metadata from the Feature Registry.  
2. Load the Features from the Offline Store. 3. Train the model.



## 2 Setup

<IPython.core.display.HTML object>

## 3 FEAST Project

Recreate the FEAST project.

---

## 4 Metadata from Feature Registry

Load the project feature metadata from the feature registry.

### 4.0.1 Verify Retrieved Metadata

```
{
  "spec": {
    "name": "customer_credit_risk",
    "description": "A project for customer credit risk"
  },
  "meta": {
    "createdTimestamp": "2025-09-18T02:27:25.053090Z",
    "lastUpdatedTimestamp": "2025-09-18T02:27:25.053090Z"
  }
}
```

```
FeatureView: customer_credit_risk_feature_view
  Feature: risk (Float32)
  Feature: purpose_business (Float32)
  Feature: purpose_car (Float32)
  Feature: purpose_domestic_appliances (Float32)
  Feature: purpose_education (Float32)
  Feature: purpose_furniture_equipment (Float32)
  Feature: purpose_radio_tv (Float32)
  Feature: purpose_repairs (Float32)
  Feature: purpose_vacation_others (Float32)
  Feature: gender_female (Float32)
  Feature: gender_male (Float32)
  Feature: housing_free (Float32)
  Feature: housing_own (Float32)
  Feature: housing_rent (Float32)
  Feature: saving_accounts_little (Float32)
  Feature: saving_accounts_moderate (Float32)
  Feature: saving_accounts_no_inf (Float32)
  Feature: saving_accounts_quite_rich (Float32)
  Feature: saving_accounts_rich (Float32)
  Feature: checking_account_little (Float32)
  Feature: checking_account_moderate (Float32)
  Feature: checking_account_no_inf (Float32)
  Feature: checking_account_rich (Float32)
  Feature: generation_student (Float32)
  Feature: generation_young (Float32)
```

Feature: generation\_adult (Float32)  
Feature: generation\_senior (Float32)  
Feature: job\_0 (Float32)  
Feature: job\_1 (Float32)  
Feature: job\_2 (Float32)  
Feature: job\_3 (Float32)  
Feature: amount\_0 (Float32)  
Feature: amount\_1 (Float32)  
Feature: amount\_2 (Float32)  
Feature: amount\_3 (Float32)  
Feature: amount\_4 (Float32)

FeatureService: customer\_credit\_risk\_feature\_service  
FeatureProjection: customer\_credit\_risk\_feature\_view  
Feature: risk (Float32)  
Feature: purpose\_business (Float32)  
Feature: purpose\_car (Float32)  
Feature: purpose\_domestic\_appliances (Float32)  
Feature: purpose\_education (Float32)  
Feature: purpose\_furniture\_equipment (Float32)  
Feature: purpose\_radio\_tv (Float32)  
Feature: purpose\_repairs (Float32)  
Feature: purpose\_vacation\_others (Float32)  
Feature: gender\_female (Float32)  
Feature: gender\_male (Float32)  
Feature: housing\_free (Float32)  
Feature: housing\_own (Float32)  
Feature: housing\_rent (Float32)  
Feature: saving\_accounts\_little (Float32)  
Feature: saving\_accounts\_moderate (Float32)  
Feature: saving\_accounts\_no\_inf (Float32)  
Feature: saving\_accounts\_quite\_rich (Float32)  
Feature: saving\_accounts\_rich (Float32)  
Feature: checking\_account\_little (Float32)  
Feature: checking\_account\_moderate (Float32)  
Feature: checking\_account\_no\_inf (Float32)  
Feature: checking\_account\_rich (Float32)  
Feature: generation\_student (Float32)  
Feature: generation\_young (Float32)  
Feature: generation\_adult (Float32)  
Feature: generation\_senior (Float32)  
Feature: job\_0 (Float32)  
Feature: job\_1 (Float32)  
Feature: job\_2 (Float32)  
Feature: job\_3 (Float32)  
Feature: amount\_0 (Float32)  
Feature: amount\_1 (Float32)  
Feature: amount\_2 (Float32)

Feature: amount\_3 (Float32)  
Feature: amount\_4 (Float32)

---

## 5 Feature Retrieval from Offline Store

- [get\\_historical\\_features](#)

This method joins historical feature data from one or more feature views to an entity dataframe by using a time travel join. Each feature view is joined to the entity dataframe using all entities configured for the respective feature view.

**Parameters**

\* **entity\_df**: a collection of rows containing all entity columns (e.g., driver\_id) on which features need to be joined, as well as a event\_timestamp column used to ensure point-in-time correctness.

**Returns**: RetrievalJob which can be used to materialize the results.

- [RetrievalJob](#)

A RetrievalJob manages the execution of a query to retrieve data from the offline store.

**Methods**

\* **to\_df**: Synchronously executes the underlying query and returns the result as a pandas dataframe. On demand transformations will be executed.

### 5.0.1 Use SQL as entity\_df

- [Example: entity SQL query for generating training data](#)

```
SELECT
    entity_id,
    event_timestamp
FROM (SELECT * FROM credit.customer_credit_risk_offline_features) AS
ENTITY_SQL
```

- [FEAST Feature Store - What is event\\_timestamp in entity\\_df parameter of Feature-Store.get\\_historical\\_features method](#)

```
entity_id          event_timestamp  risk  purpose_business  purpose_car
purpose_domestic_appliances  purpose_education  purpose_furniture_equipment
purpose_radio_tv  purpose_repairs  purpose_vacation_others  gender_female
gender_male  housing_free  housing_own  housing_rent  saving_accounts_little
saving_accounts_moderate  saving_accounts_no_inf  saving_accounts_quite_rich
saving_accounts_rich  checking_account_little  checking_account_moderate
checking_account_no_inf  checking_account_rich  generation_student
generation_young  generation_adult  generation_senior  job_0  job_1  job_2
job_3  amount_0  amount_1  amount_2  amount_3  amount_4
```

```

0          1 2025-09-18 12:27:23.123026 0.0          0.0          0.0
0.0          0.0          0.0          0.0          1.0          0.0
0.0          0.0          0.0          0.0          1.0          0.0
1.0          0.0          0.0          0.0          0.0          0.0
1.0          0.0          0.0          0.0          0.0          1.0
0.0          0.0          0.0          0.0          0.0          0.0
0.0          0.0          0.0          1.0          0.0          0.0          1.0          0.0          1.0
0.0          0.0          0.0          0.0          0.0          0.0          1.0          0.0          1.0
1          2 2025-09-18 12:27:23.123026 1.0          0.0          0.0
0.0          0.0          0.0          0.0          0.0          1.0
0.0          0.0          0.0          1.0          0.0          0.0
1.0          0.0          0.0          1.0          0.0          0.0
0.0          0.0          0.0          0.0          0.0          0.0
1.0          0.0          0.0          0.0          0.0          1.0
0.0          0.0          0.0          0.0          0.0          0.0
1.0          0.0          0.0          0.0          0.0          0.0          1.0          0.0          0.0
2          3 2025-09-18 12:27:23.123026 0.0          0.0          0.0
0.0          1.0          0.0          0.0          0.0          0.0
0.0          0.0          0.0          0.0          1.0          0.0
1.0          0.0          0.0          1.0          0.0          0.0
0.0          0.0          0.0          0.0          0.0          0.0
0.0          1.0          0.0          0.0          0.0          0.0
0.0          1.0          0.0          0.0          0.0          0.0
0.0          1.0          0.0          0.0          1.0          0.0          0.0          0.0          1.0
0.0          0.0          0.0          0.0

```

## 5.1 Online Store

Need to run the materialization command to get columns from the online store.

### 5.1.1 Materialise

Materializing 1 feature views from 2025-09-18

12:27:23+00:00 to 2025-09-18 12:27:45+00:00 into the  
sqlite online store.

customer\_credit\_risk\_feature\_view:

### 5.1.2 Issue

[FeatureStore.get\\_online\\_features\(\) throws FeatureViewNotFoundException #5606](#)

```

FeatureViewNotFoundException: Feature view cus-
tomer_credit_risk_customer_credit_risk_feature_view does not exist in project
customer_credit_risk

```

```

entity_id  gender_female  checking_account_moderate
purpose_domestic_appliances  saving_accounts_moderate  amount_4  purpose_repairs
amount_0  generation_senior  housing_rent  checking_account_little  risk

```

```

purpose_business amount_3 generation_student housing_own
saving_accounts_little checking_account_rich checking_account_no_inf
purpose_education generation_adult saving_accounts_quite_rich
generation_young amount_1 housing_free job_1 amount_2
purpose_vacation_others purpose_radio_tv gender_male saving_accounts_no_inf
purpose_car purpose_furniture_equipment saving_accounts_rich job_0 job_2
job_3
0          1          0.0          0.0          0.0          1.0
0.0          0.0          0.0          0.0          0.0          0.0
1.0          0.0          1.0          0.0          0.0          0.0
0.0          1.0          0.0          0.0          0.0          0.0
0.0          0.0          0.0          0.0          0.0          0.0
0.0          0.0          0.0          0.0          0.0          0.0
1.0          1.0          1.0          0.0          0.0          0.0
0.0          0.0          0.0          1.0          0.0          0.0
1          2          1.0          1.0          0.0          0.0
0.0          0.0          0.0          0.0          0.0          0.0
0.0          0.0          0.0          1.0          0.0          0.0
1.0          1.0          1.0          0.0          0.0          0.0
0.0          0.0          0.0          0.0          0.0          0.0
0.0          1.0          0.0          0.0          0.0          0.0
1.0          0.0          0.0          0.0          0.0          0.0
0.0          0.0          0.0          1.0          0.0          0.0

```

---

## 6 Model Training - Feature Consumption

```

purpose_business purpose_car purpose_domestic_appliances purpose_education
purpose_furniture_equipment purpose_radio_tv purpose_repairs
purpose_vacation_others gender_female gender_male housing_free housing_own
housing_rent saving_accounts_little saving_accounts_moderate
saving_accounts_no_inf saving_accounts_quite_rich saving_accounts_rich
checking_account_little checking_account_moderate checking_account_no_inf
checking_account_rich generation_student generation_young generation_adult
generation_senior job_0 job_1 job_2 job_3 amount_0 amount_1 amount_2
amount_3 amount_4
0          0.0          0.0          0.0          0.0          0.0
0.0          1.0          0.0          0.0          0.0          0.0
1.0          0.0          1.0          0.0          0.0          0.0
0.0          1.0          0.0          0.0          0.0          0.0
1.0          0.0          0.0          0.0          0.0          0.0
0.0          0.0          0.0          1.0          0.0          0.0
0.0          1.0          0.0          0.0          0.0          0.0
1          0.0          0.0          0.0          0.0          0.0
0.0          1.0          0.0          0.0          0.0          1.0
0.0          0.0          1.0          0.0          1.0          0.0

```

0.0		0.0		0.0		0.0
0.0		1.0		0.0		0.0
1.0		0.0		0.0	0.0	0.0
0.0	0.0	1.0	0.0	0.0	0.0	1.0
2		0.0	0.0		0.0	1.0
0.0		0.0		0.0	0.0	0.0
1.0	0.0		1.0	0.0		1.0
0.0		0.0		0.0		0.0
0.0		0.0		1.0		0.0
0.0		0.0		1.0	0.0	0.0
0.0	1.0	0.0	0.0	0.0	0.0	1.0
0.0		0.0		0.0		0.0

## 6.1 Algorithm Comparision

XGB: 0.486110 (0.104101)

Gaussian Naive Bayes: 0.525474 (0.082965)

Logistic Regression: 0.599749 (0.080013)

## 7 Train Models

## 7.1 Logistic Regression

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
GridSearchCV(cv=5,
             estimator=LogisticRegression(class_weight='balanced',
                                           max_iter=1000, random_state=42,
                                           solver='liblinear'),
             n_jobs=-1,
             param_grid={'C': [0.01, 0.1, 1, 10, 100],
                         'class_weight': ['balanced', None],
                         'penalty': ['l2']},
             scoring='f1', verbose=2)
```

Logistic regression best score: 0.5894372219310449

```
{
  "C": 0.01,
  "class_weight": "balanced",
  "penalty": "l2"
}
```

Best parameters: {'C': 0.01, 'class\_weight': 'balanced', 'penalty': 'l2'}

Best ROC: 0.5894372219310449

### Logistic Regression Results:

### Confusion Matrix

[[115 63]]

[ 21 51]]

Accuracy: 0.6640

Recall: 0.7083


Precision: 0.4474

F1: 0.5484

AUC: 0.7365

classification report

	precision	recall	f1-score	support
0.0	0.85	0.65	0.73	178
1.0	0.45	0.71	0.55	72
accuracy			0.66	250
macro avg	0.65	0.68	0.64	250
weighted avg	0.73	0.66	0.68	250



NOTE: SKLearn Confusion Matrix format:

```
[[TN FP]
 [FN TP]]
```



## 7.2 GNB

- [Scikit-Learn GridSearchCV GaussianNB](#)

0.6946666666666668

```
GridSearchCV(cv=KFold(n_splits=10, random_state=7, shuffle=True),
             estimator=Pipeline(steps=[('feature_union',
                                       FeatureUnion(transformer_list=[('pca',
PCA(n_components=2)),
('select_best',
SelectKBest(k=6))])),
             ('logistic', GaussianNB())]),
            n_jobs=-1,
            param_grid={'feature_union__pca__n_components': [1, 2, 3],
                        'feature_union__select_best__k': [4, 6, 8],
                        'logistic__var_smoothing': [1e-09, 1e-08, 1e-07,
                                                    1e-06]},
            scoring='f1')
```

GNB best score: 0.5516661003209277

```
{
  "feature_union__pca__n_components": 3,
  "feature_union__select_best__k": 4,
  "logistic__var_smoothing": 1e-09
}
```

Gaussian NB Results:

Confusion Matrix

```
[[128  50]
```

```
 [ 27  45]]
```

Accuracy: 0.6920

Recall: 0.6250

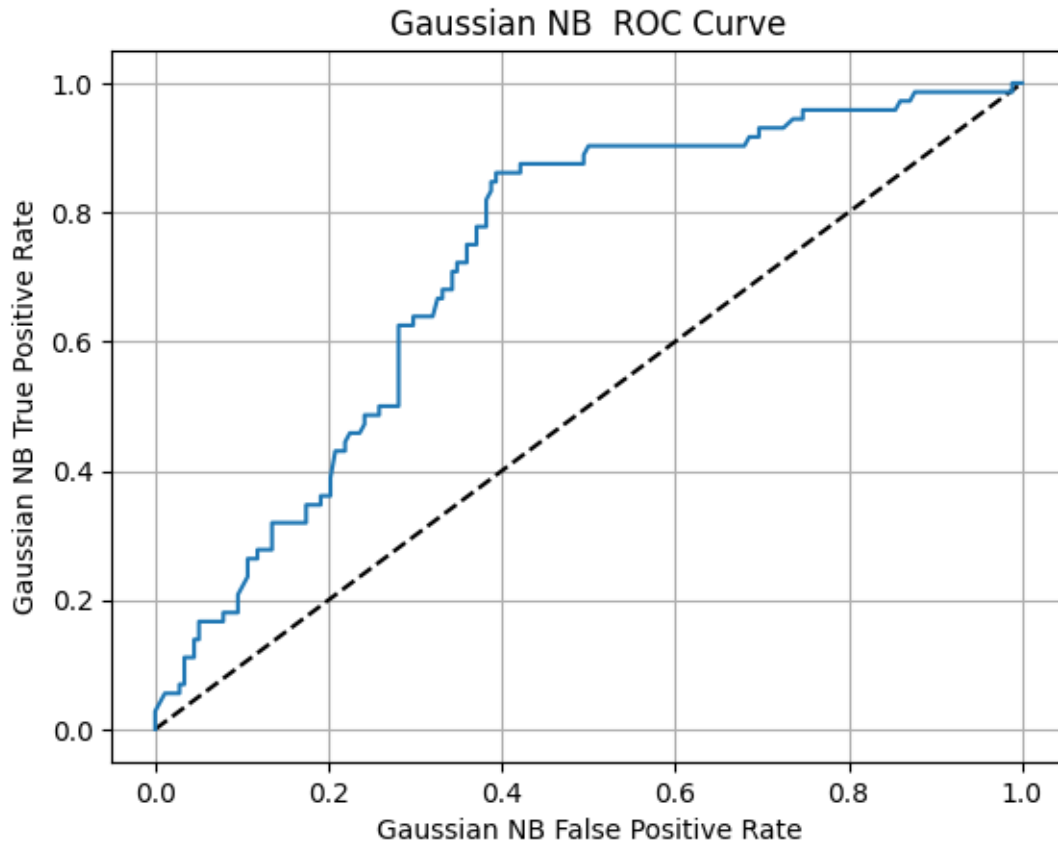
Precision: 0.4737

F1: 0.5389

AUC: 0.7239

classification report

	precision	recall	f1-score	support
0.0	0.83	0.72	0.77	178
1.0	0.47	0.62	0.54	72
accuracy			0.69	250
macro avg	0.65	0.67	0.65	250
weighted avg	0.72	0.69	0.70	250



### 7.3 XGB

```
GridSearchCV(cv=5,
             estimator=XGBClassifier(base_score=None, booster=None,
                                     callbacks=None, colsample_bylevel=None,
                                     colsample_bynode=None,
                                     colsample_bytree=None, device=None,
                                     early_stopping_rounds=None,
                                     enable_categorical=False,
                                     eval_metric='auc', feature_types=None,
                                     feature_weights=None, gamma=None,
                                     grow_policy=None, importance_type=None,
                                     interaction_constraints=None,
                                     max_delta_step=None, max_depth=None,
                                     max_leaves=None, min_child_weight=None,
                                     missing=nan, monotone_constraints=None,
                                     multi_strategy=None, n_estimators=None,
                                     n_jobs=-1, num_parallel_tree=None, ...),
             param_grid={'gamma': [0.1, 0.2]},
```

```

        'learning_rate': [0.005, 0.01, 0.015],
        'max_depth': [3, 4], 'min_child_weight': [8, 12],
        'n_estimators': [100, 200, 300],
        'scale_pos_weight': [2.3333333333333335]}},
    scoring='f1')

grid_xb.best_score_: 0.5782563728585046
grid_xb.best_params_ {
  "gamma": 0.1,
  "learning_rate": 0.01,
  "max_depth": 3,
  "min_child_weight": 12,
  "n_estimators": 200,
  "scale_pos_weight": 2.3333333333333335
}

```

XGBoost Results:

Confusion Matrix

```
[[110  68]
```

```
 [ 17  55]]
```

Accuracy: 0.6600

Recall: 0.7639


Precision: 0.4472

F1: 0.5641

AUC: 0.7231

classification report

	precision	recall	f1-score	support
0.0	0.87	0.62	0.72	178
1.0	0.45	0.76	0.56	72
accuracy			0.66	250
macro avg	0.66	0.69	0.64	250
weighted avg	0.75	0.66	0.68	250



	feature	importance
20	checking_account_no_inf	0.283878
21	checking_account_rich	0.064194
14	saving_accounts_moderate	0.058702
29	job_3	0.058189
18	checking_account_little	0.048575
22	generation_student	0.046834
15	saving_accounts_no_inf	0.046445
24	generation_adult	0.044408
11	housing_own	0.043040
30	amount_0	0.038106